

Data Analyst Interview Preparation

Detailed Interview Questions & Answers — Power BI, DAX, SQL, Data Warehousing, Python/API, PySpark & Microsoft Fabric

This expanded guide is based on the 64-question reference PDF provided for interview preparation. The original reference organizes the material into Power BI, DAX, SQL, Data Modeling, ETL/Data Warehouse, Python/API, Microsoft Fabric, and communication questions.

How to use this guide

Do not memorize every answer word-for-word. Understand the concept, learn the example, and practice saying the interview-ready version naturally. For technical questions, be prepared to explain both the definition and a simple example.

Power BI Basics

1. What is Power BI used for?

Power BI is a Microsoft Business Intelligence and data visualization tool. It is used to connect data from different sources, clean and transform it, create a data model, perform calculations using DAX, and present business information through interactive reports and dashboards.

Example: A company can connect Excel or SQL sales data to Power BI and build a report showing total sales, profit, monthly trends, products, customers, and regions.

Interview-ready: "Power BI is a Microsoft BI tool used to connect, transform, model, and visualize data. I can use Power Query for cleaning, DAX for calculations, create interactive reports, and publish them to Power BI Service."

Remember: Connect → Clean → Model → DAX → Visualize → Publish.

2. What steps do you follow to build a Power BI report?

I first connect Power BI to the required data source. Next, I use Power Query to clean and transform the data—for example, handling nulls, changing data types, removing duplicates, merging or appending tables. Then I build the data model and create relationships between tables. After that, I create DAX measures, design visuals, and add filters and slicers. Finally, I validate the numbers and publish the report to Power BI Service.

Interview-ready: "My usual flow is Connect → Clean → Model → DAX → Visualize → Validate → Publish."

3. Dashboard vs Report

A report is designed for detailed analysis and can contain multiple pages with many visuals, filters, and slicers. A dashboard is generally a high-level monitoring view made up of tiles and is useful for quickly checking important KPIs.

Example: A sales report may have separate pages for Overview, Products, Customers, and Regions. A management dashboard may show only Total Sales, Profit, Orders, and a few key charts.

Interview-ready: "A report is mainly for detailed analysis and can have multiple pages, while a dashboard is mainly for high-level monitoring and quick insights."

4. What is a dataset / semantic model?

A semantic model is the data model behind a Power BI report. It contains tables, relationships, measures, columns, and business logic used by the report. It provides a structured business view of the underlying data.

Example: A sales semantic model may contain Sales, Product, Customer, and Date tables, relationships between them, and measures such as Total Sales and YTD Sales.

Interview-ready: "A semantic model is the structured model behind a Power BI report. It contains data tables, relationships, measures, and business logic that visuals use for analysis."

5. What is Power Query used for?

Power Query is used to connect to data sources and perform data preparation before the data is loaded into the model. Typical tasks include removing duplicates, handling nulls, changing data types, splitting or merging columns, filtering rows, merging tables, and appending tables.

Example: If an Excel file contains duplicate customers and dates stored as text, Power Query can remove duplicates and convert the date column to a proper date type.

Interview-ready: "I use Power Query mainly for data extraction, cleaning, and transformation before modeling."

6. Measures vs Calculated Columns

A calculated column is calculated row by row and its values are stored in the model. A measure is a calculation evaluated when it is used in a visual and responds to the current filter context.

Example: A calculated column could calculate $\text{SalesAmount} = \text{Quantity} \times \text{UnitPrice}$ for every sales row. A measure could calculate $\text{Total Sales} = \text{SUM}(\text{Sales}[\text{SalesAmount}])$.

Interview-ready: "I use a calculated column when I need a value at row level, while I use a measure for dynamic aggregations and KPIs that should respond to filters."

7. Import vs DirectQuery

In Import mode, data is loaded into Power BI's in-memory model, so reports are generally fast. Data must be refreshed to reflect source changes. In DirectQuery, data remains in the source system and Power BI sends queries to the source when users interact with the report.

Example: Import can be useful when the dataset can be loaded and refreshed periodically. DirectQuery can be useful when the organization needs to query data that should remain in the source.

Remember: Import = data comes to Power BI. DirectQuery = Power BI queries the data source.

8. What is a slicer?

A slicer is an interactive visual filter. Users select one or more values in a slicer, and connected visuals are filtered accordingly.

Example: If a Year slicer contains 2024 and 2025, selecting 2025 can make the sales charts display only 2025 data.

Interview-ready: "A slicer is a visual filter that allows users to interactively select values and filter the report visuals."

9. How do you publish a report?

I create and validate the report in Power BI Desktop, save it, and select Publish. I then choose the required workspace in Power BI Service. After publishing, I check the report and semantic model, configure refresh if required, and make sure users have the appropriate access.

Interview-ready: "I build and validate the report in Desktop, publish it to the required Power BI Service workspace, and then configure refresh and access as needed."

10. What is refresh? Why does it fail?

Refresh updates the data in the Power BI model from its source. Refresh can fail because of database connectivity problems, incorrect credentials, gateway issues, unavailable sources, changed table or column names, or transformation errors.

A good troubleshooting approach is: check the refresh error, verify the source connection and credentials, check gateway status where applicable, confirm that source objects still exist, and review Power Query steps.

Interview-ready: "Refresh means updating the model with the latest source data. If it fails, I check connectivity, credentials, gateway, source changes, and transformation errors."

11. Name 3 visuals you used and why

Bar charts are useful for comparing categories, such as sales by product. Line charts are useful for showing trends over time, such as monthly sales. Cards are useful for displaying important KPIs such as Total Sales, Profit, or Customer Count.

Interview-ready: "I commonly use bar charts for category comparisons, line charts for time trends, and cards for important KPIs."

DAX Interview Questions

12. SUM vs CALCULATE

SUM adds the values in a column. CALCULATE evaluates an expression after modifying or applying filter context.

Example: Total Sales = SUM(Sales[SalesAmount]). A CALCULATE expression can change the filters before calculating the result.

Interview-ready: "SUM performs a straightforward aggregation, while CALCULATE is used when I need to modify filter context for a calculation."

13. Row Context vs Filter Context

Row context means DAX is evaluating one row at a time, which is common in calculated columns and iterator functions. Filter context is the set of filters affecting a calculation, coming from slicers, visuals, relationships, or DAX.

Example: Selecting 2025 in a Year slicer creates filter context so a Total Sales measure calculates for the selected year.

Interview-ready: "Row context evaluates data row by row, while filter context determines which data is included in a calculation."

14. SUM vs SUMX

SUM adds an existing numeric column. SUMX is an iterator: it evaluates an expression row by row and then adds the results.

Example:

SUM(Sales[SalesAmount])

SUMX(Sales, Sales[Quantity] * Sales[UnitPrice])

Use SUM when the required value already exists in a column. Use SUMX when you need a row-level calculation before aggregation.

15. What is context transition?

Context transition occurs when CALCULATE converts an existing row context into filter context. This is important when CALCULATE is used inside an iterator or in situations where a current row needs to influence a measure calculation.

Interview-ready: "Context transition is the process where CALCULATE turns the current row context into filter context."

16. What is YTD Sales?

YTD means Year To Date. It calculates a value from the beginning of the current year up to the current date.

DAX example:

YTD Sales = TOTALYTD([Total Sales], 'Date'[Date])

Example: If the current report date is March 31, YTD Sales represents sales from January 1 through March 31.

Interview-ready: "YTD calculates the cumulative value from the start of the year through the current date."

17. What is MoM Growth %?

MoM means Month over Month. It compares the current month's value with the previous month's value.

Typical calculation:

MoM Growth % = DIVIDE([Total Sales] - [Previous Month Sales], [Previous Month Sales])

Example: If January sales are 100 and February sales are 120, growth is 20%.

Interview-ready: "MoM growth measures how much the current month increased or decreased compared with the previous month."

18. What is a Running Total?

A running total is an accumulated value up to the current date or point in an ordered sequence.

Example: If daily sales are 10, 20, and 30, the running totals are 10, 30, and 60.

It is useful for cumulative sales, cumulative orders, or cumulative profit analysis.

19. What is Rolling 7 Days?

A rolling 7-day calculation calculates a value over a moving seven-day window. For a given day, it normally includes that day and the previous six days.

Example: The rolling value for July 10 uses July 4 through July 10.

It is useful for smoothing daily fluctuations and analyzing recent performance.

20. ALL / ALLEXCEPT / ALLSELECTED

ALL removes filters from the specified table or columns. ALLEXCEPT removes filters except those on specified columns.

ALLSELECTED generally keeps the user's broader selections while changing the context created by the current visual.

Interview-ready: "I use ALL when I need to ignore filters, ALLEXCEPT when I want to preserve selected grouping columns, and ALLSELECTED when I need a calculation that respects the user's selections while handling visual context."

21. How do you debug wrong numbers in DAX?

I do not immediately change the formula. First, I check the source data and confirm the expected number. Then I check table relationships, filter and slicer behavior, data types, and the DAX formula. I also create small test measures to isolate each part of the calculation.

A useful approach is: Source → Relationships → Filters → Data Types → DAX → Visual.

Interview-ready: "I compare the Power BI result with the source, then isolate the issue by checking relationships, filters, data types, and smaller test measures."

SQL Interview

22. Write a query using JOIN

JOIN combines related data from multiple tables. For example, employees can be joined to departments using department ID.

SQL:

```
SELECT e.name, d.department
```

```
FROM employees e
```

```
INNER JOIN departments d
```

```
ON e.dept_id = d.dept_id;
```

INNER JOIN returns rows where a matching department exists.

Interview-ready: "I use JOIN when related information is stored in different tables and I need to combine it using a common key."

23. GROUP BY + HAVING

GROUP BY creates groups so aggregate functions such as COUNT, SUM, and AVG can be calculated for each group.

HAVING filters the groups after aggregation.

SQL:

```
SELECT department, COUNT(*) AS employee_count
```

```
FROM employees
```

```
GROUP BY department
```

```
HAVING COUNT(*) > 2;
```

WHERE filters individual rows before grouping; HAVING filters grouped results.

24. How do you find duplicate records?

Group the records by the column or combination of columns that should be unique and use HAVING COUNT(*) > 1.

SQL:

```
SELECT name, COUNT(*)
```

```
FROM employees
```

```
GROUP BY name
```

```
HAVING COUNT(*) > 1;
```

This identifies names that appear more than once. In real projects, I would use the appropriate business key rather than assuming name alone defines a duplicate.

25. What is ROW_NUMBER?

ROW_NUMBER is a window function that assigns a unique sequential number to rows based on the ORDER BY inside the window.

SQL:

```
SELECT name, salary,
```

```
ROW_NUMBER() OVER(ORDER BY salary DESC) AS rn
```

```
FROM employees;
```

The highest salary gets row number 1. If two employees have the same salary, ROW_NUMBER still gives them different numbers unless additional ordering is used.

26. Primary Key vs Foreign Key

A primary key uniquely identifies a record in its table and cannot contain NULL values. A foreign key is a column that references a key in another table and is used to establish relationships.

Example: department_id can be the primary key in Departments and a foreign key in Employees.

Interview-ready: "Primary key identifies a row uniquely; foreign key connects one table to another."

27. WHERE vs HAVING

WHERE filters rows before GROUP BY and aggregation. HAVING filters groups after GROUP BY.

Example:

```
SELECT department, COUNT(*)  
FROM employees  
WHERE salary > 50000  
GROUP BY department  
HAVING COUNT(*) > 2;
```

Here WHERE removes low-salary rows first, while HAVING keeps only departments with more than two remaining employees.

28. INNER JOIN vs LEFT JOIN

INNER JOIN returns only rows with a match in both tables. LEFT JOIN returns every row from the left table and matching rows from the right table; if there is no match, right-side columns are NULL.

Interview-ready: "I use INNER JOIN when I only need matching records and LEFT JOIN when all records from the left table must be retained."

29. Find top 5 products by revenue

```
SQL:  
SELECT product_id, SUM(revenue) AS total_revenue  
FROM sales  
GROUP BY product_id  
ORDER BY total_revenue DESC  
LIMIT 5;
```

First we group sales by product, then calculate total revenue, sort from highest to lowest, and return the first five products. Note: LIMIT syntax can differ across database systems.

30. Count customers per city

```
SQL:  
SELECT city, COUNT(*) AS customer_count  
FROM customers  
GROUP BY city;
```

GROUP BY creates one group for each city and COUNT calculates the number of records in each group.

31. What is a subquery?

A subquery is a query inside another SQL query. It can be used to provide a value or dataset to the outer query.

Example:

```
SELECT name, salary  
FROM employees  
WHERE salary > (SELECT AVG(salary) FROM employees);
```

The inner query calculates average salary, and the outer query finds employees earning more than that average.

32. What is indexing?

An index is a database structure that can help locate rows faster, similar to an index in a book. Indexes can improve read/query performance, but they require storage and can add overhead to INSERT, UPDATE, and DELETE operations.

Interview-ready: "I use indexes thoughtfully on columns frequently used for filtering, joining, or ordering, while considering their write and storage cost."

33. Normalization vs Denormalization

Normalization organizes data into related tables to reduce duplication and improve consistency. Denormalization intentionally combines or duplicates some data to simplify queries or improve reporting performance.

In transactional systems, normalization is common. In analytical systems, denormalized structures such as star schemas are often preferred.

Interview-ready: "Normalization reduces redundancy; denormalization can make analytical queries simpler and faster."

34. How do you handle NULLs?

First I understand why the value is NULL and what the business rule says. Depending on the requirement, I can filter NULLs with IS NULL/IS NOT NULL, replace values using COALESCE, or leave them as NULL when NULL has business meaning.

Example:

```
SELECT COALESCE(phone, 'Not Available') AS phone  
FROM customers;
```

35. RANK vs DENSE_RANK

Both are ranking window functions, but they handle ties differently. If salaries are 70000, 70000, 60000, RANK returns 1, 1, 3 because it skips a rank after the tie. DENSE_RANK returns 1, 1, 2 because it does not skip a rank.

Use RANK when competition-style ranking with gaps is expected; use DENSE_RANK when consecutive ranking values are required.

36. UNION vs UNION ALL

UNION combines compatible result sets and removes duplicate rows. UNION ALL combines result sets but keeps duplicates. Because UNION has to perform duplicate elimination, UNION ALL is often preferable when duplicates are valid or already controlled.

Interview-ready: "UNION removes duplicates; UNION ALL keeps them."

Data Modeling

37. What is a Fact Table?

A fact table stores measurable business events or transactions. Typical columns include foreign keys to dimensions and numeric measures such as quantity, sales, cost, or profit.

Example: Sales Fact can contain DateKey, ProductKey, CustomerKey, Quantity, SalesAmount, and Cost.
The grain of a fact table is important: it defines exactly what one row represents.

38. What is a Dimension Table?

A dimension table contains descriptive attributes used to analyze facts. Common dimensions include Customer, Product, Date, Employee, and Location.

Example: Product Dimension may contain ProductKey, ProductName, Category, Brand, and Subcategory.

Dimensions answer questions such as “who?”, “what?”, “where?”, and “when?” while facts usually contain measurable events.

39. What is Star Schema?

A star schema is a data model where a central fact table connects directly to multiple dimension tables. It is called a star because the relationships visually resemble a star.

Example: Sales Fact in the center connected to Date, Product, Customer, and Store dimensions.

40. Why do we prefer Star Schema?

Star schema is relatively simple to understand, supports clear relationships, works well for analytical queries, and is commonly used with Power BI. It separates measurable events from descriptive attributes and makes business analysis easier.

Interview-ready: “I prefer a star schema for reporting because it is simple, has clear relationships, and works well for analytical queries and Power BI models.”

41. What is a Date Table?

A date table is a dedicated table containing a continuous set of dates and useful attributes such as Year, Month, Month Number, Quarter, and Day.

It is important for time intelligence calculations such as YTD, MTD, QTD, previous year, and MoM.

Example columns: Date, Year, Quarter, Month, Month Number, Day, Day of Week.

ETL + Data Warehouse

42. What is ETL?

ETL stands for Extract, Transform, Load. Extract means collecting data from sources such as databases, files, or APIs. Transform means cleaning, validating, joining, standardizing, and applying business rules. Load means writing the prepared data into the target data warehouse.

Interview-ready: "ETL is a process where data is extracted from sources, transformed according to business requirements, and loaded into a data warehouse."

43. Full Load vs Incremental Load

A full load processes all required data each time. An incremental load processes only new or changed records.

Example: If a table contains 10 million records and only 10,000 changed today, incremental loading can process those changed records instead of all 10 million.

Incremental loading can reduce processing time and resource usage, but it requires a reliable method for identifying new or changed records.

44. What is Data Quality?

Data quality means ensuring that data is accurate, complete, consistent, valid, and reliable enough for its intended use. Typical checks include NULL values, duplicates, invalid data types, primary and foreign key integrity, valid ranges, unexpected values, and source-to-target record counts.

Interview-ready: "I treat data quality as a set of checks that confirm the data is accurate, complete, consistent, valid, and reliable."

45. What is a Surrogate Key?

A surrogate key is an artificial, system-generated identifier used to uniquely identify a record, especially in dimension tables.

Example: A Customer Dimension might use Customer_SK = 1001 as its warehouse key, while the source system has Customer_ID = C458.

Surrogate keys are particularly useful when tracking historical changes in dimensions.

46. What is SCD?

SCD stands for Slowly Changing Dimension. It is a data warehousing technique for managing changes in dimension attributes over time.

Example: If a customer's city changes, the warehouse needs a rule defining whether to overwrite the old city or preserve historical information.

47. SCD Type 1 vs Type 2

SCD Type 1 overwrites the old value, so previous history is not retained. It is useful when only the current value matters. SCD Type 2 preserves history by creating a new dimension record for the changed version, commonly with effective dates and a current-record indicator.

Example: If a customer moves from Delhi to Mumbai, Type 1 replaces Delhi with Mumbai. Type 2 keeps the Delhi version and creates a new Mumbai version.

Interview-ready: "Type 1 overwrites; Type 2 maintains historical versions."

Python / API

48. How do you read a CSV in Python/Pandas?

Pandas provides `read_csv()` to load a CSV file into a DataFrame.

Code:

```
import pandas as pd
df = pd.read_csv("sales.csv")
```

After loading, I can inspect it with `df.head()`, `df.info()`, and `df.describe()` and then clean or analyze the data.

Interview-ready: "I use `pandas.read_csv()` to load CSV data into a DataFrame, then inspect and transform it as required."

49. How do you handle missing values?

First I identify missing values and understand why they are missing. Then I choose an approach based on the business requirement: remove records, replace values with 0 when zero is meaningful, or use statistics such as mean or median where appropriate.

Example:

```
df.isna().sum()
df["salary"] = df["salary"].fillna(df["salary"].median())
```

I would not blindly replace every NULL with zero because NULL and zero can have different meanings.

50. List vs Dictionary

A list is an ordered collection of values accessed mainly by position/index. A dictionary stores data as key-value pairs and is accessed using keys.

Example:

```
names = ["Amit", "Riya", "Neha"]
employee = {"name": "Amit", "salary": 50000}
```

A list is useful for sequences; a dictionary is useful when each value has a meaningful key.

51. What is groupby in Pandas?

`groupby` is used to divide data into groups based on one or more columns and then perform calculations such as sum, count, mean, minimum, or maximum.

Example:

```
df.groupby("department")["salary"].sum()
```

This calculates total salary for each department.

Interview-ready: "I use `groupby` to perform grouped aggregations and summarize data by business dimensions."

52. What is an API?

API stands for Application Programming Interface. It provides a defined way for applications or systems to communicate.

Example: A Python application can call a REST API to retrieve customer or weather data in JSON format.

Interview-ready: "An API is an interface that allows different software systems to communicate and exchange data or functionality."

53. GET vs POST

GET is generally used to retrieve data from a server. POST is generally used to send data to a server, often to create a new resource.

Example:

GET /students → retrieve students.

POST /students → send student information to create a student.

Interview-ready: "GET is mainly for retrieving data, while POST is mainly for sending data or creating a resource."

54. What is Pagination?

Pagination means dividing a large API response into smaller pages rather than returning every record in one response.

Example: An API may return 100 customers per page and provide page=1, page=2, and so on.

Pagination reduces response size and can improve application performance and usability.

55. What is a Rate Limit?

A rate limit controls how many requests a client can make to an API during a specified time period.

Example: An API might allow only a certain number of requests per minute.

Rate limits protect services from excessive traffic and help maintain fair access for users.

56. Pandas vs PySpark

Pandas is generally suitable for smaller or medium-sized datasets that can be processed on one machine and fit into available memory. PySpark is designed for distributed processing and is useful for very large datasets across multiple machines.

Interview-ready: "I would choose Pandas for data that can be handled efficiently on one machine and PySpark when distributed processing is needed for very large data."

57. Spark Transformation vs Action

A Spark transformation defines an operation or produces a new DataFrame/RDD lineage, while an action triggers execution and returns or materializes a result.

Examples of transformations include filter(), select(), and groupBy(). Examples of actions include count(), show(), and collect().

Interview-ready: "Transformations define what Spark should do; actions trigger execution and produce a result."

Microsoft Fabric

58. What is Microsoft Fabric?

Microsoft Fabric is an end-to-end analytics platform from Microsoft. It brings together data engineering, data integration, data warehousing, data science, real-time analytics, and Power BI in one platform.

Power BI mainly focuses on analytics and reporting, while Fabric covers a broader data and analytics lifecycle.

Interview-ready: "Microsoft Fabric is a unified analytics platform that brings multiple data and analytics workloads together."

59. What is OneLake?

OneLake is the centralized data lake for Microsoft Fabric. It provides a common place to store and access organizational data across Fabric workloads.

Interview-ready: "OneLake is Fabric's centralized data lake that provides a single data foundation across Fabric workloads."

60. Lakehouse vs Data Warehouse in Fabric

A Lakehouse combines characteristics of a data lake and data warehouse and is useful for flexible storage, data engineering, and analytics workloads. A Fabric Warehouse is designed primarily for structured relational data and SQL-based analytics.

Interview-ready: "I would consider a Lakehouse when flexible data storage and engineering are important, while Fabric Warehouse is a strong choice for structured relational and SQL analytics."

61. What is a Fabric Pipeline?

A Fabric pipeline is used to automate and orchestrate data movement and processing tasks. A pipeline can coordinate activities such as extracting data, transforming it, loading it into a destination, and running tasks on a schedule.

Interview-ready: "A Fabric pipeline automates and orchestrates data workflows so extraction, transformation, and loading can run in a controlled sequence."

Communication / Project Questions

62. Explain your project in 2 minutes

"I worked on a data analytics project where I collected data from sources such as SQL and Excel. I used SQL for data extraction and transformation and Power Query for data cleaning. I created a star schema with fact and dimension tables and developed DAX measures such as Total Sales, YTD Sales, and Growth %. Then I created an interactive Power BI report using cards, charts, and slicers. Finally, I published the report to Power BI Service and configured refresh. The main objective was to provide business users with an easy way to analyze performance and make better decisions."

Tip: While speaking, explain your actual project details where possible rather than memorizing this answer word-for-word.

63. What if the report shows wrong numbers?

First, I would not immediately change the DAX formula. I would compare the Power BI result with the source data or SQL result. Then I would check the relationships between tables, filters and slicers, data types, duplicate rows, and the DAX calculation. I would create smaller test measures if necessary to identify exactly where the number changes.

Interview-ready: "I troubleshoot from source to report: Source → Relationships → Filters → Data Types → DAX → Visual, and compare results with SQL where possible."

64. What challenges did you face?

A good example is data quality issues such as NULL values, duplicate records, inconsistent values, or incorrect data types. These can be handled during transformation using SQL and Power Query. Another common challenge is incorrect numbers in visuals, which can be investigated by checking relationships, filter context, and DAX measures.

Interview-ready: "One challenge was inconsistent data quality. I handled it during transformation using SQL and Power Query. When visuals showed unexpected numbers, I checked relationships, filters, and DAX calculations to identify the root cause."

Priority for Interview Preparation

Priority 1 — Must Know

Power BI; Power Query; Report vs Dashboard; Import vs DirectQuery; Measure vs Calculated Column; SUM vs SUMX; CALCULATE; Row Context vs Filter Context; YTD; SQL JOIN; WHERE vs HAVING; GROUP BY; Primary Key / Foreign Key; ROW_NUMBER; RANK vs DENSE_RANK; Fact vs Dimension; Star Schema; ETL; Full Load vs Incremental Load; SCD Type 1 vs Type 2.

Priority 2

ALL / ALLEXCEPT / ALLSELECTED; Context Transition; Running Total; MoM Growth; Duplicate records; Subquery; Indexing; Normalization; Data quality checks; Surrogate key.

Priority 3 — JD/Fabric

Microsoft Fabric; OneLake; Lakehouse; Fabric Warehouse; Fabric Pipeline; Pandas; API; GET vs POST; Pagination; PySpark.

Final Interview Tip

A useful project flow to remember is: Source → SQL → ETL → Model → DAX → Power BI → Publish. The reference material specifically recommends understanding concepts rather than memorizing answers word-for-word.